

COMP 3647

Human-AI Interaction Design

Topic 15

Ethics in AI

Dr. Swaroop Panda,
Prof. Effie Law

Roadmap

- Tech Ethics
- Explainable & Responsible AI
- ML Fairness
- Bias & Discrimination
- Human Autonomy & Privacy
- Governance & Legal Aspects

What is Ethics

- Aristotle's Ethics
- Normative Ethics
 - Virtue Ethics
 - Deontology
 - Consequentialism

Tech Ethics

- The Trolley Problem




Image by McGeddon (CC BY-SA 4.0)

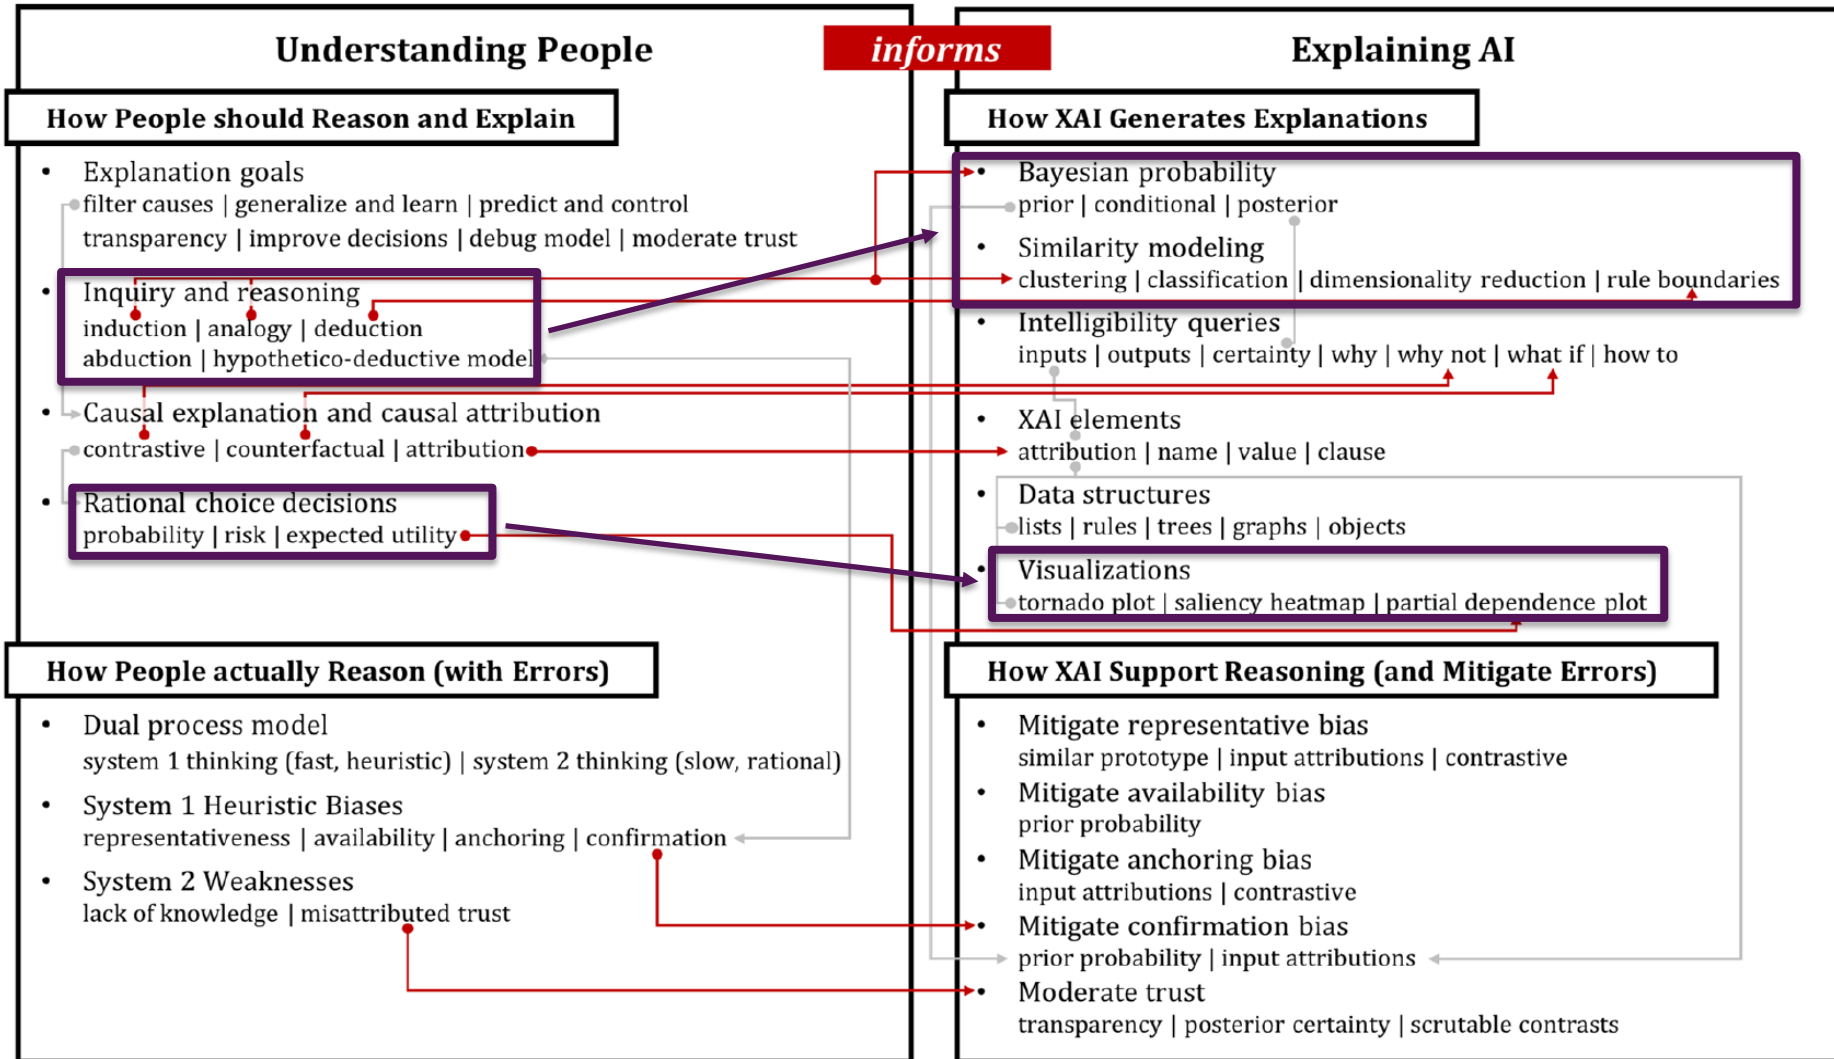
Tech Ethics

You have two options:

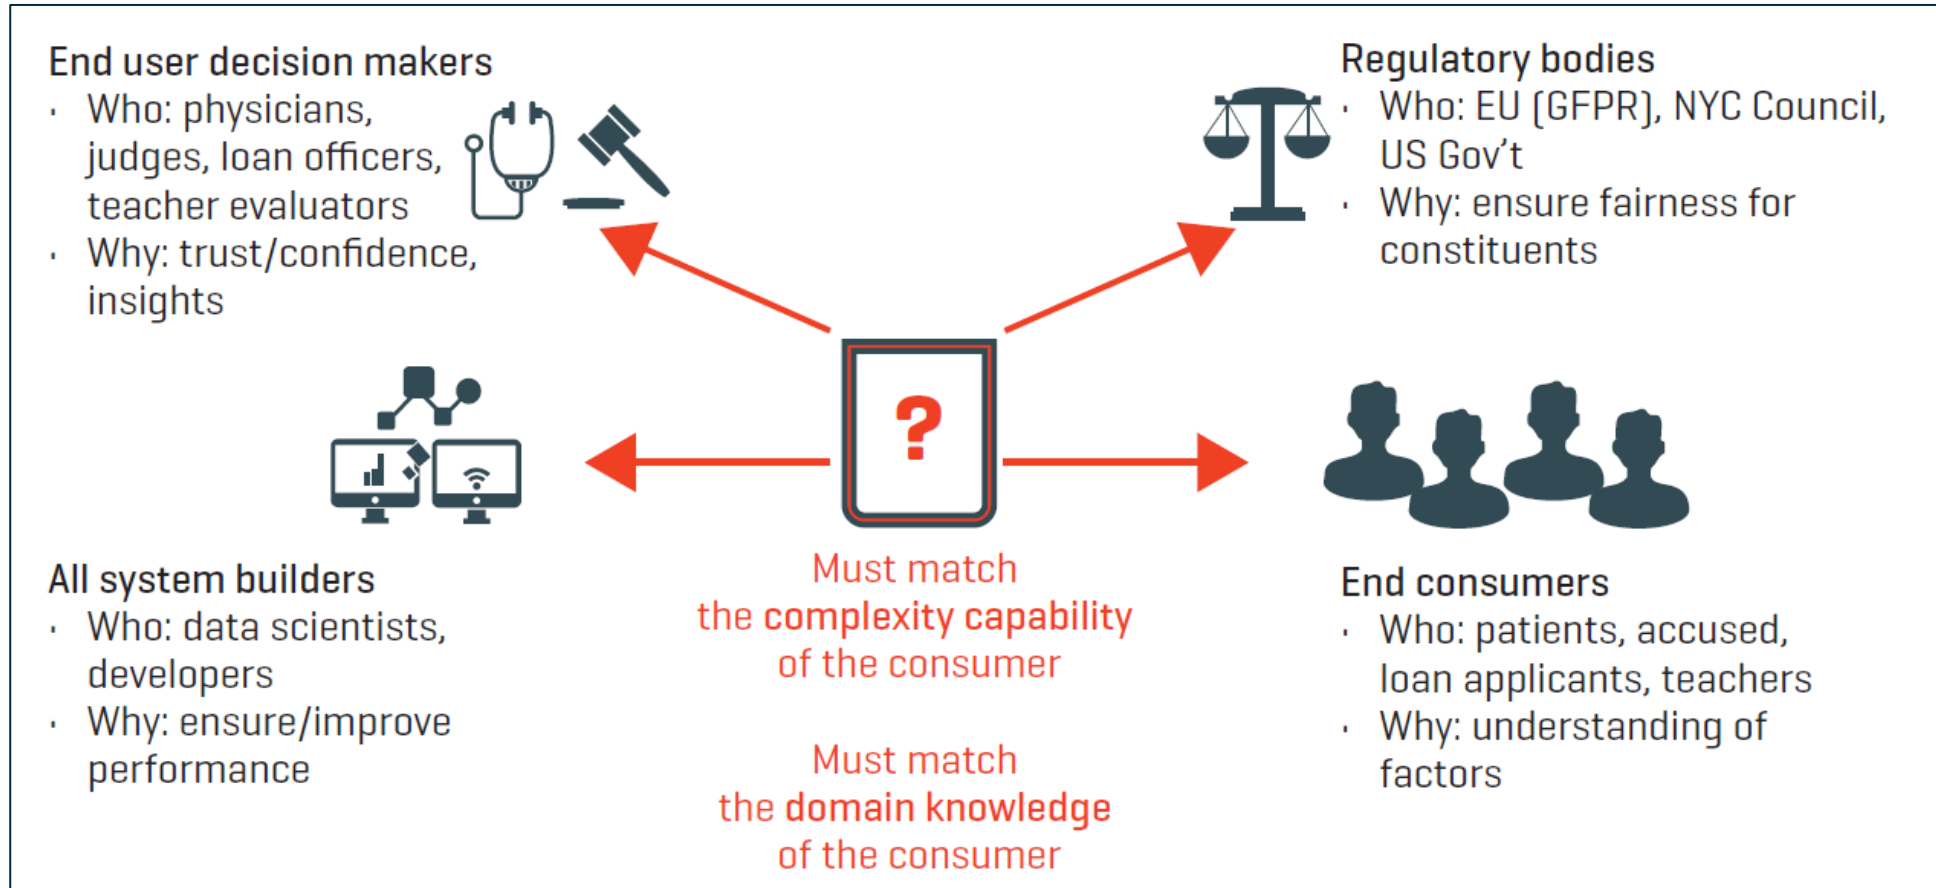
1. **Do nothing**, and the trolley kills the five people on the main track.
2. **Pull the lever**, diverting the trolley onto the side track where it will kill one person.

Which is the most **ethical choice**?

Explainable AI



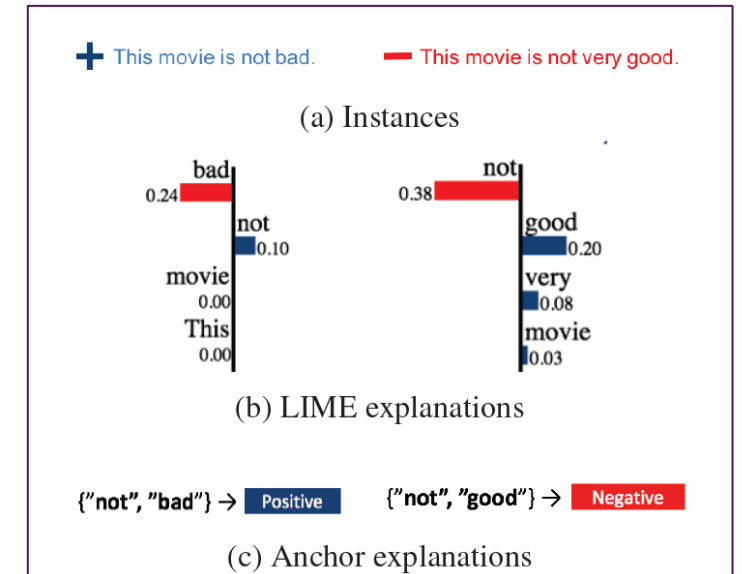
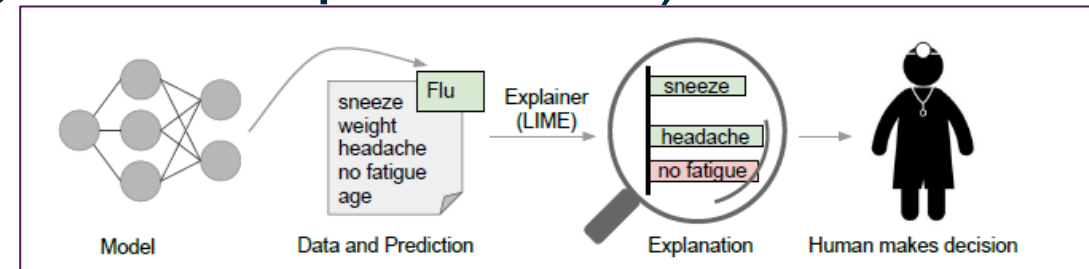
Explainable AI



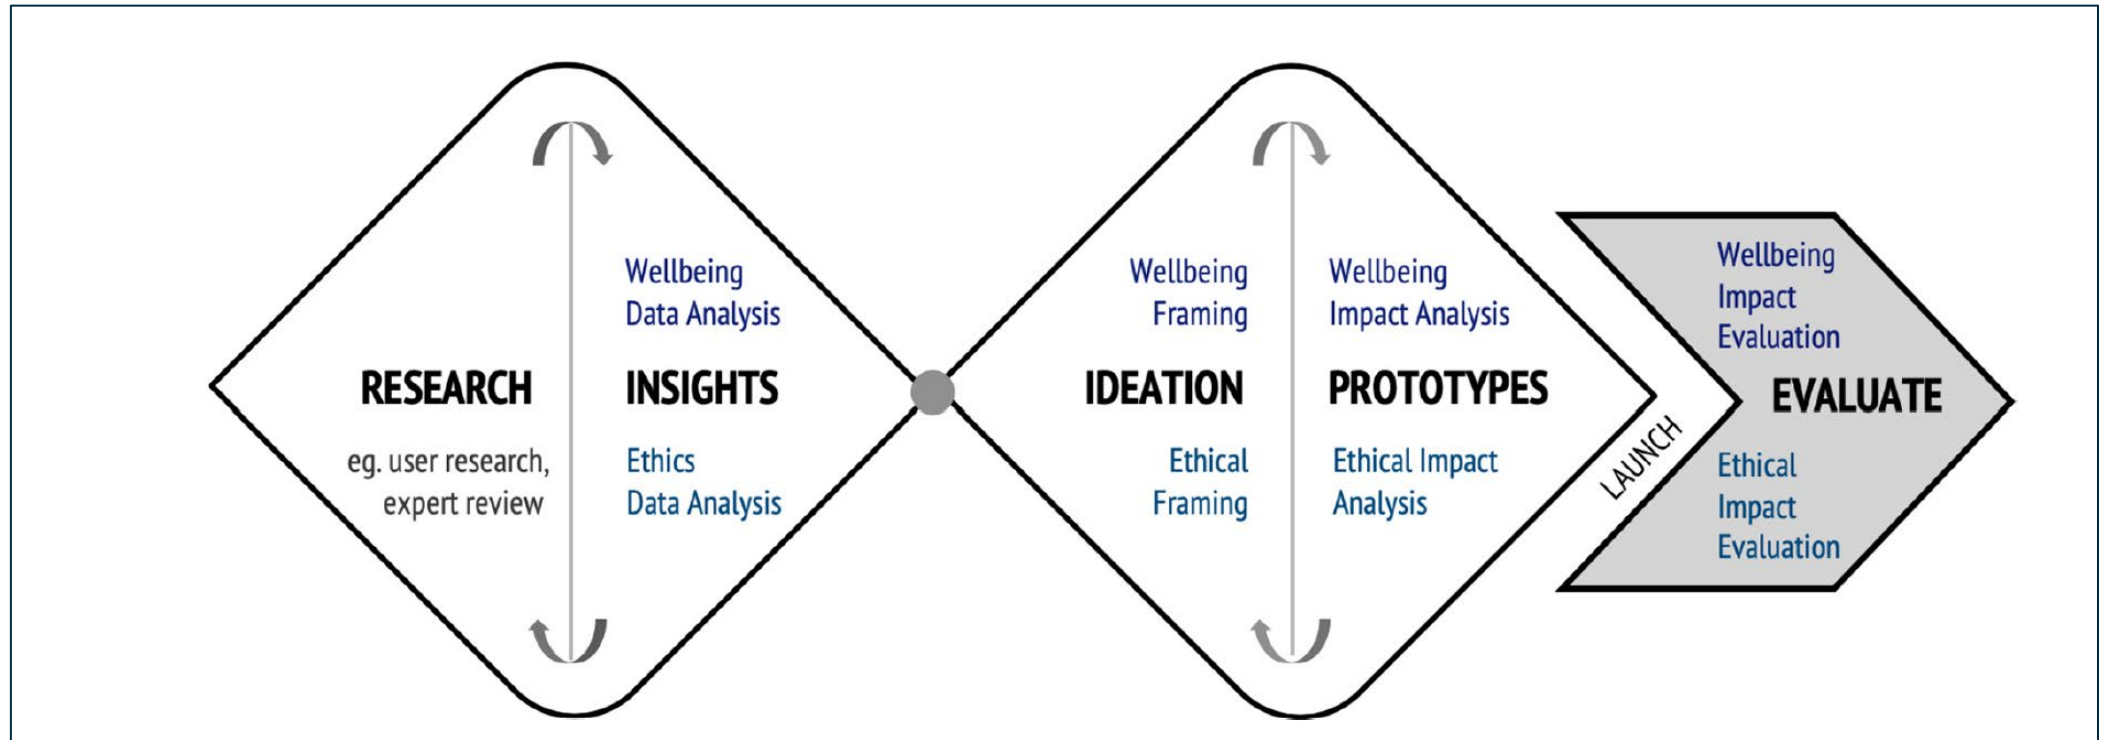
Explainable AI

Some XAI Methods,

1. LIME (Local Interpretable Model Agnostic Explanations)
2. Anchors
3. Layer-wise Relevance Propagation
4. Deep Taylor Decomposition (DTD)
5. Others



Responsible AI



ML Fairness

Confusion Matrix

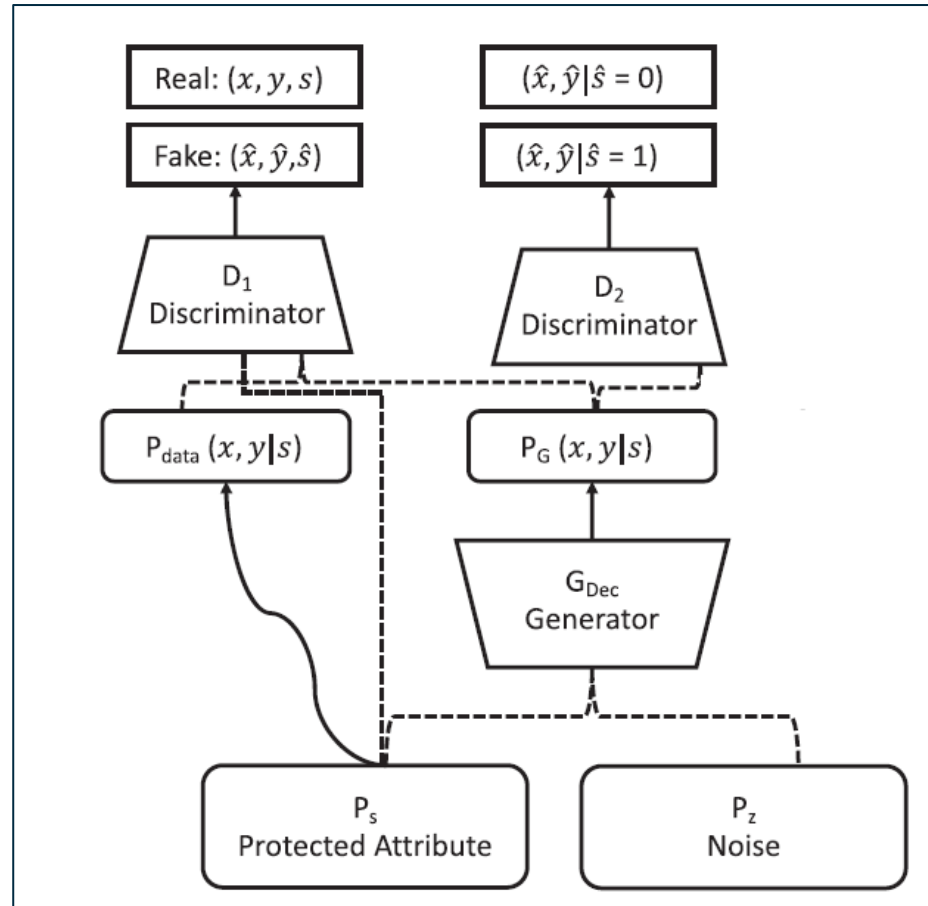
	Actual Positive $Y = 1$	Actual Negative $Y = 0$		
Predicted Positive $\hat{Y} = 1$	TP (True Positive)	FP (False Positive) Type I error	PPV = $\frac{TP}{TP+FP}$ Positive Predictive Value Precision PV+ Target Population Error	FDR = $\frac{FP}{TP+FP}$ False Discovery Rate Target Population Error
Predicted Negative $\hat{Y} = 0$	FN (False Negative) Type II error	TN (True Negative)	FOR = $\frac{FN}{FN+TN}$ False Omission Rate Success Predictive Error	NPV = $\frac{TN}{FN+TN}$ Negative Predictive Value PV-
	TPR = $\frac{TP}{TP+FN}$ True Positive Rate Sensitivity Recall	FPR = $\frac{FP}{FP+TN}$ False Positive Rate Model Error	OA = $\frac{TP+TN}{TP+FP+TN+FN}$ Overall Accuracy	BR = $\frac{TP+FN}{TP+FP+TN+FN}$ Base Rate Prevalence (p)
	FNR = $\frac{FN}{TP+FN}$ False Negative Rate Model Error	TNR = $\frac{TN}{FP+TN}$ True Negative Rate Specificity		

ML Fairness

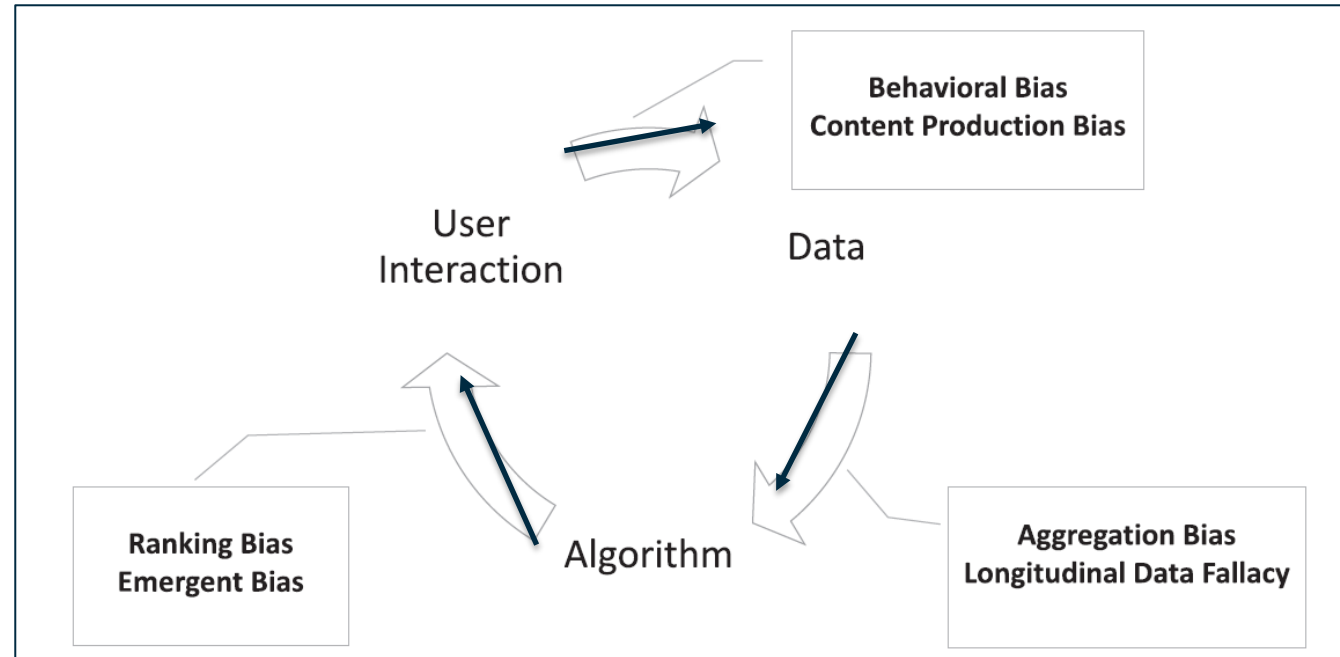
Definitions,

- Equalised odds
- Equal Opportunity
- Demographic Parity
- Fairness through awareness
- Test-fairness or calibration
- Others ...

ML Fairness - FairGAN



ML Fairness



Bias & Discrimination

Types of Bias

- Measurement Bias
- Omitted Variable Bias
- Representation Bias
- Aggregation Bias
 - Simpson's Paradox.
 - Modifiable Areal Unit Problem

Bias & Discrimination

Types of Bias

- Algorithmic Bias
- User Interaction Bias
- Popularity Bias
- Emergent Bias
- Evaluation Bias
- Population Bias
- Historical Bias & Others

Bias & Discrimination

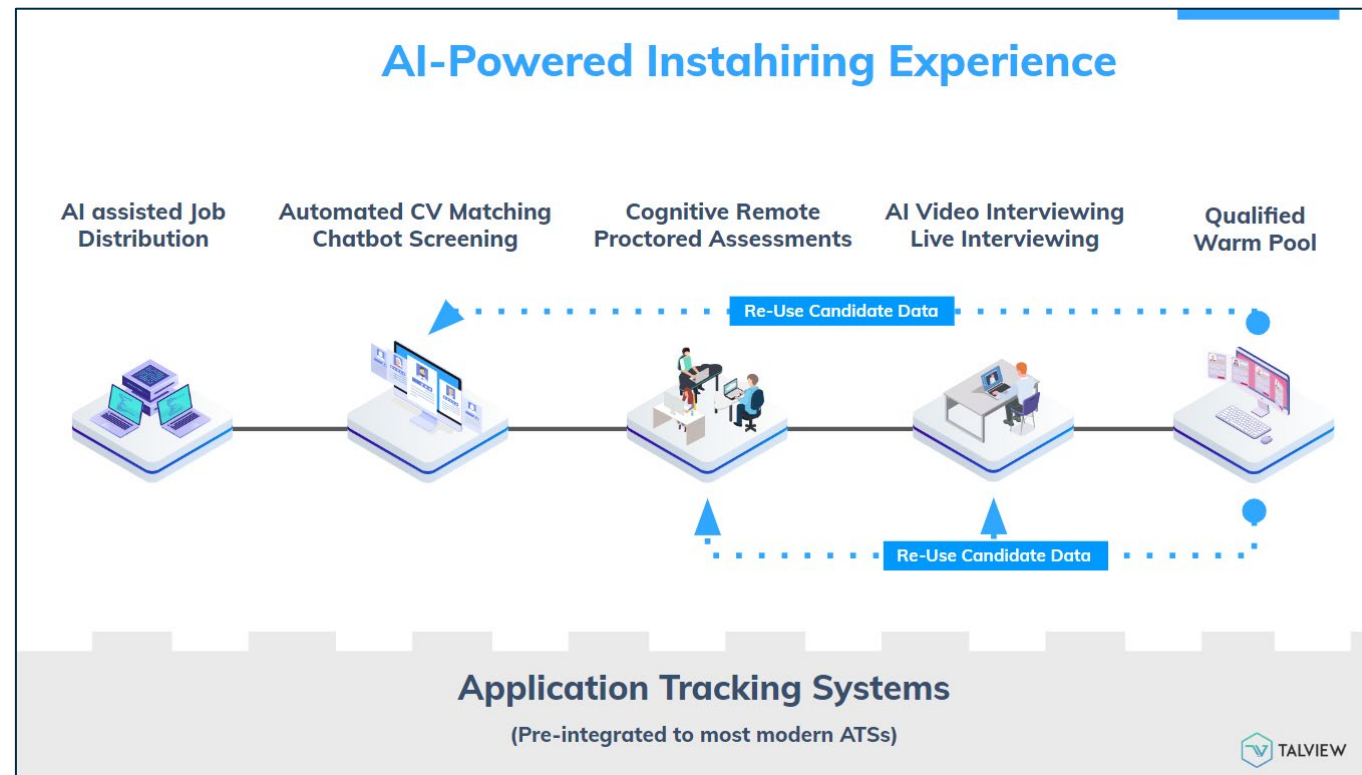
Discrimination vs Bias

Types of Discrimination,

- Systemic Discrimination
- Statistical Discrimination

Human Autonomy

- Replacing Human Labor



Human Autonomy

- Autonomous Weapons



Human Autonomy

- Replacing Humans



Privacy

- Is Siri/Alexa hearing us?



Privacy

Surveillance



Privacy


- Facial Recognition



Legal Aspects

- Who should regulate AI?

**From a 'race to AI' to a 'race to AI regulation':
regulatory competition for artificial intelligence**

Nathalie A. Smuha 

Faculty of Law, KU Leuven, Leuven, Belgium

Legal Aspects

- Who should regulate AI?



Legal Aspects

Equality



The assumption is that **everyone benefits from the same supports**. This is equal treatment.

Equity



Everyone gets the supports they need (this is the concept of "affirmative action"), thus producing equity.

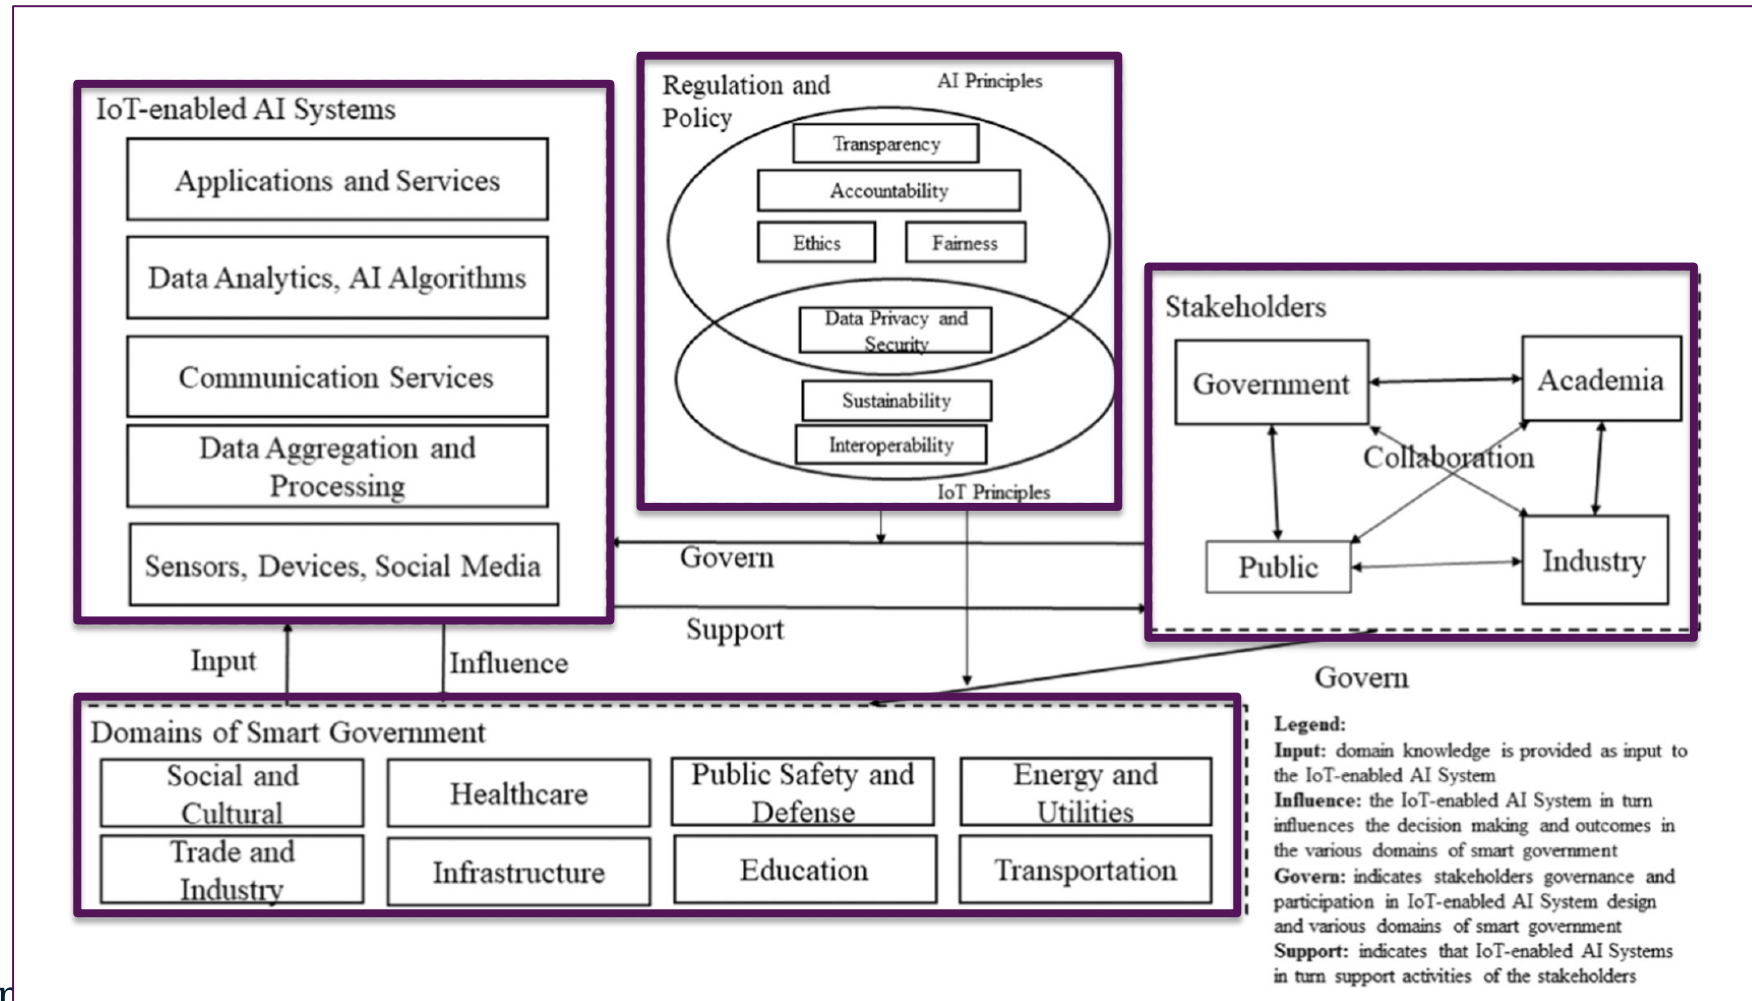
Justice



All 3 can see the game without supports or accommodations because **the cause(s) of the inequity was addressed**. The systemic barrier has been removed.


Governance

- Smart Governance



Governance

Freedom at Work: Understanding, Alienation, and the AI-Driven Workplace

Kate Vredenburg 

Department of Philosophy, Logic, and Scientific Method, The London School of Economics, London, United Kingdom
Email: K.Vredenburg@lse.ac.uk

Governance

Designing AI with Rights, Consciousness, Self-Respect, and Freedom

Eric Schwitzgebel, with Mara Garza

Governance

- AI to oppress

Organizing workers and machine learning tools for a less
oppressive workplace

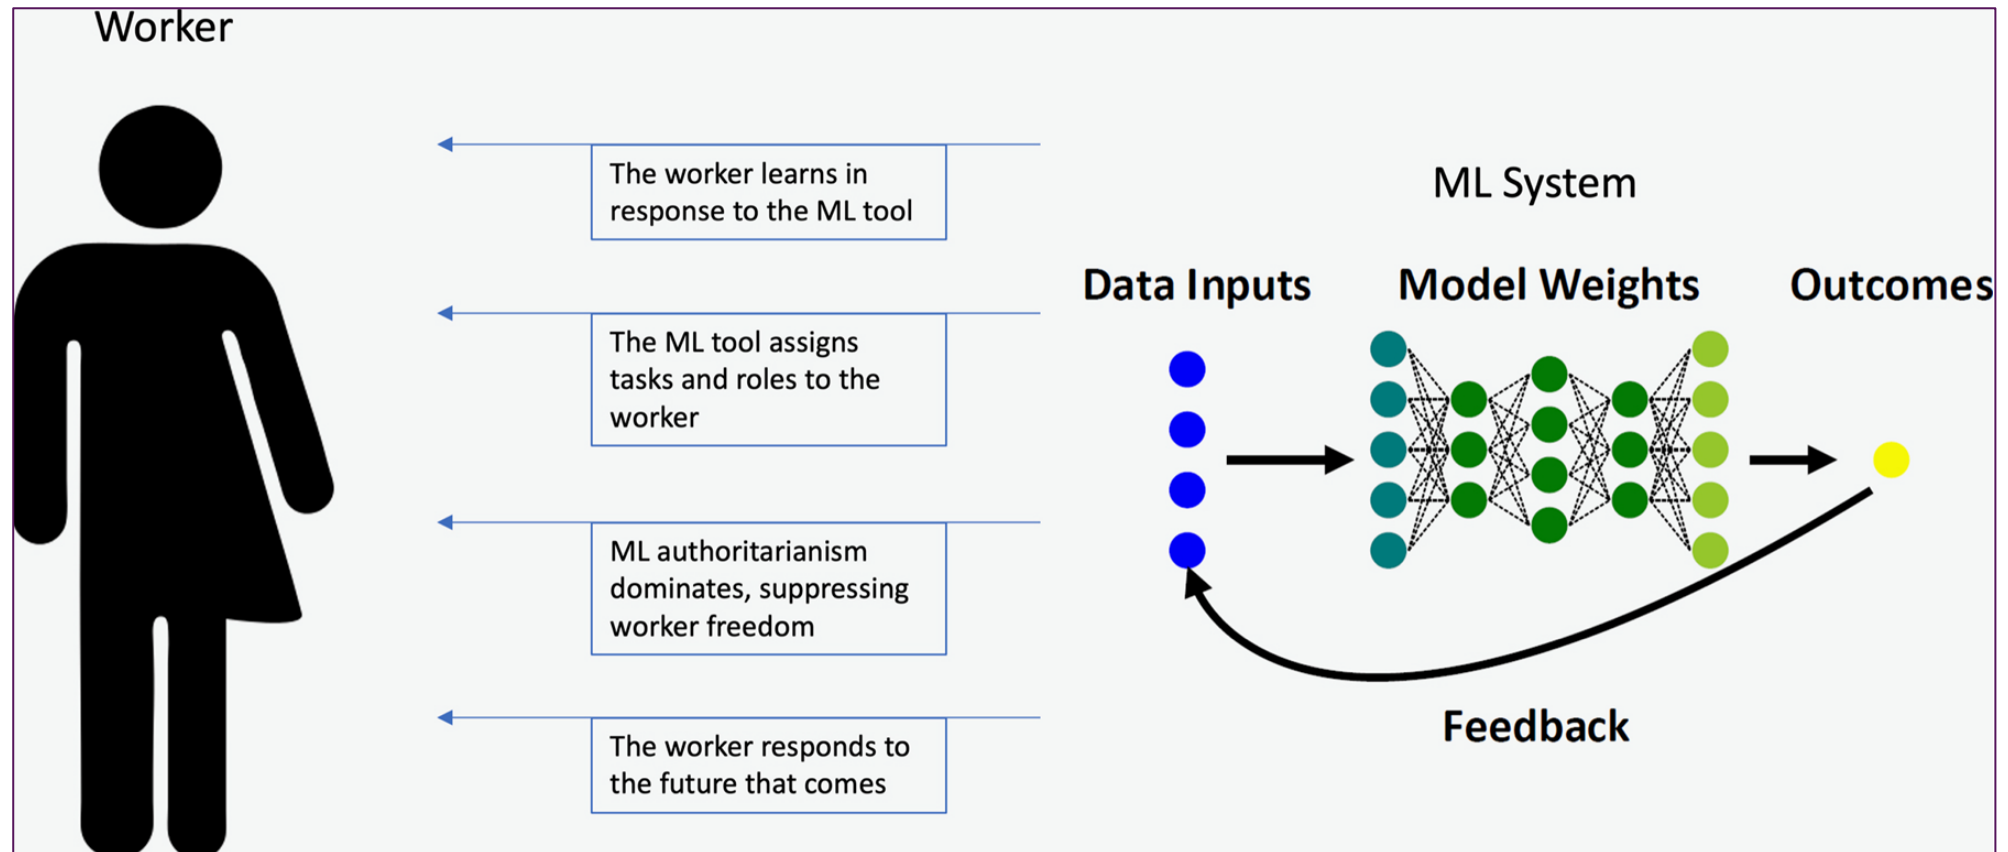
Amber Grace Young^{a,*}, Ann Majchrzak^b, Gerald C. Kane^c

^a Sam M. Walton College of Business, University of Arkansas, 220 N McIlroy Ave #301, Fayetteville, AR, 72701, USA

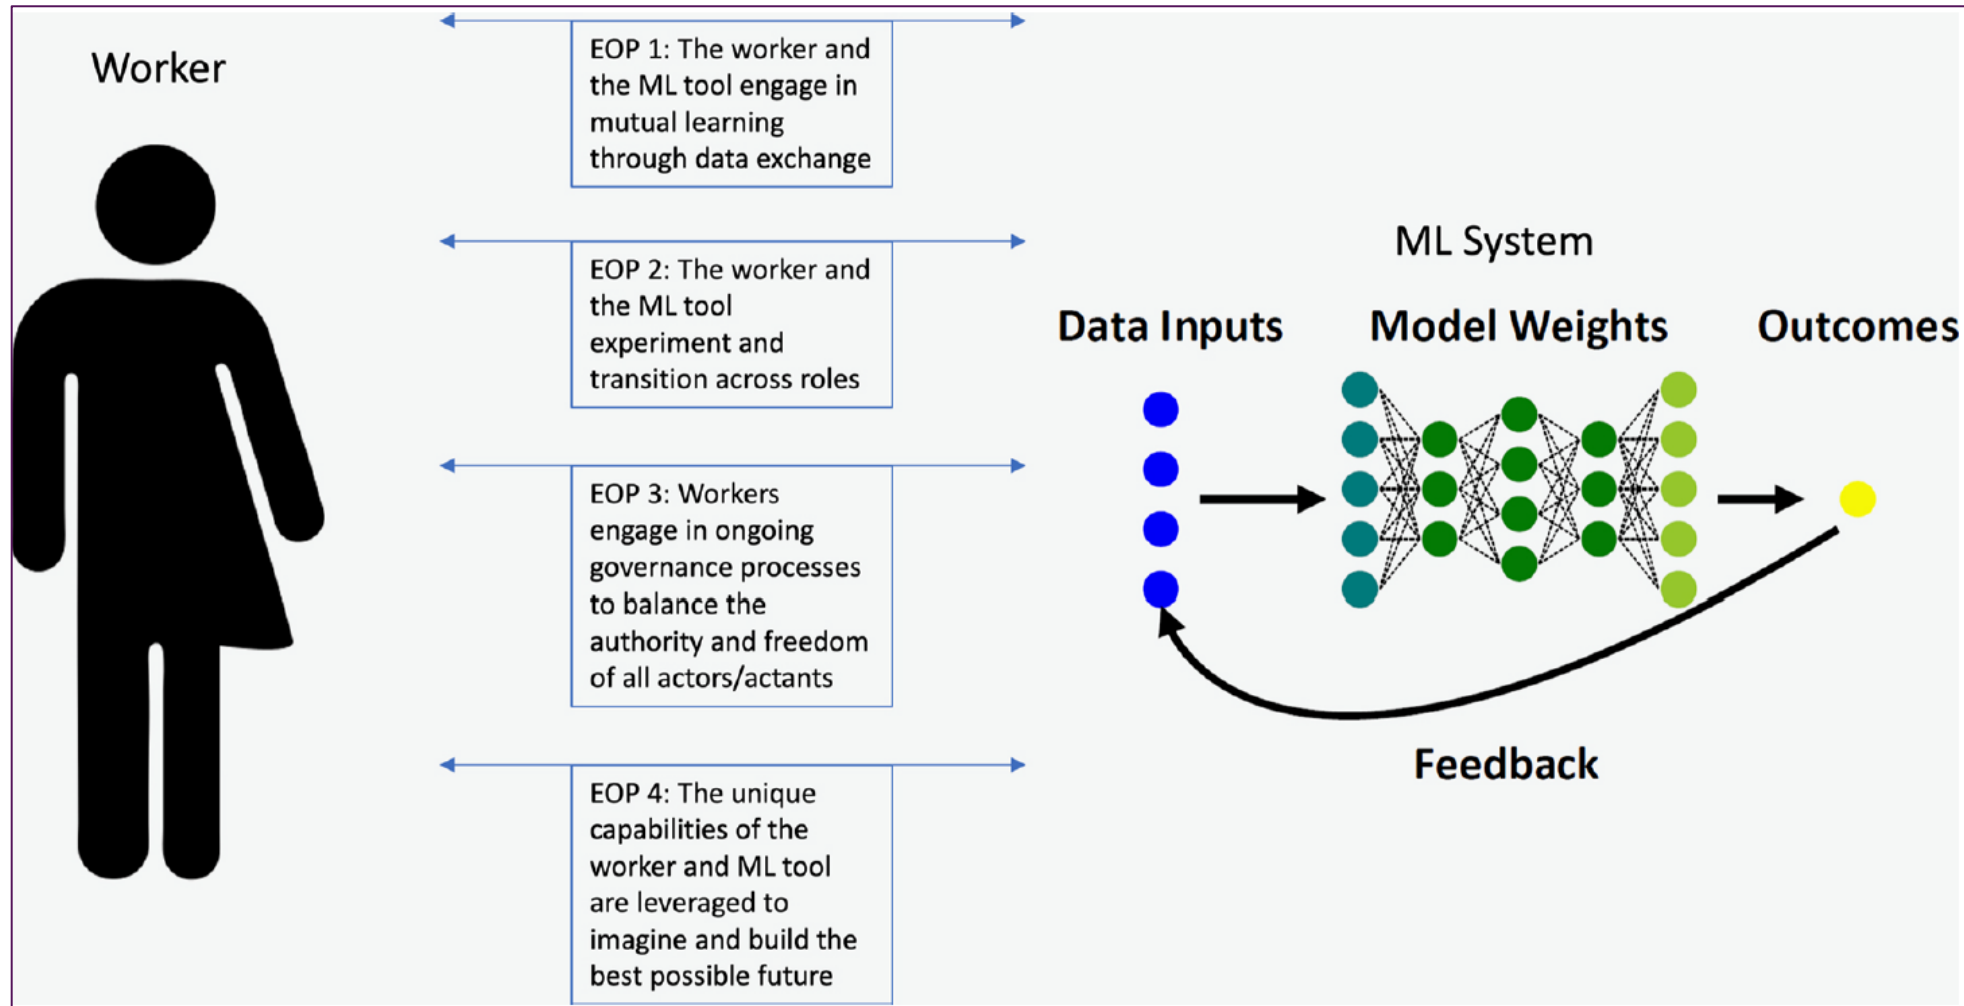
^b Marshall School of Business, University of Southern California, 3670 Trousdale Pkwy, Los Angeles, CA, 90089, USA

^c Carroll School of Management, Boston College, 140 Commonwealth Avenue, Chestnut Hill, MA, 02467, USA

Governance



Governance



Wrapping Up

- Ethics in Technology
- Explainable & Responsible AI Methods
- Fairness in AI
- Bias & Discrimination
- Human Autonomy & Privacy
- Governance & Legal Aspects

Thanks

Any Questions?